

Direct Neural Machine Translation with Task-level Mixture of Experts models

Isidora Chara Tourni *

Boston University
isidora@bu.edu

Subhajit Naskar

Google
snaskar@google.com

Abstract

Direct neural machine translation (direct NMT) is a type of NMT system that translates text between two non-English languages. Direct NMT systems often face limitations due to the scarcity of parallel data between non-English language pairs. Several approaches have been proposed to address this limitation, such as multilingual NMT and pivot NMT (translation between two languages via English). Task-level Mixture of expert models (Task-level MoE), an inference-efficient variation of Transformer-based models, has shown promising NMT performance for a large number of language pairs. In Task-level MoE, different language groups can use different routing strategies to optimize cross-lingual learning and inference speed. In this work, we examine Task-level MoE’s applicability in direct NMT and propose a series of high-performing training and evaluation configurations, through which Task-level MoE-based direct NMT systems outperform bilingual and pivot-based models for a large number of low and high-resource direct pairs, and translation directions. Our Task-level MoE model with 16 experts outperforms bilingual NMT, Pivot NMT models for 7 language pairs, while pivot-based models still performed better in 9 pairs and directions.

1 Introduction

NMT (Sutskever et al., 2014; Bahdanau et al., 2014; Firat et al., 2016a; Johnson et al., 2017; Zoph and Knight, 2016; Bahdanau et al., 2014) has made remarkable progress in recent years, particularly with the advent of the Transformer architecture (Vaswani et al., 2017), which has shown impressive results in bilingual and multilingual translation tasks (Aharoni et al., 2019; Hassan et al., 2018; Tang et al., 2021). Large Multilingual NMT models in particular have shown noteworthy performance

and can serve translations between any languages, even for language pairs not seen in training, for which though NMT quality is often still low. The focus so far has primarily been on English-centric NMT, as there is an abundance of English-centric MT data (Tiedemann, 2018) but few parallel data for non-English (low-resource) language pairs. Hence the need for improving methods for Direct NMT, i.e. NMT between non-English languages is evident (Yang et al., 2021; Zhang et al., 2020).

Direct NMT methods, when no xx-yy data is available (where xx represents the source language code, and yy represents the target language code), can be categorized into Zero-shot and Zero-resource approaches. In Zero-shot NMT (Johnson et al., 2017; Gu et al., 2019; Wang et al., 2021; Philip et al., 2020), translation occurs as follows: either through a bridging language, typically English, where models are trained on En-yy and xx-En datasets without ever being exposed to xx-yy data, while evaluation takes place on the few existing xx-yy pairs; or evaluating large multilingual models on xx-yy pairs, with neither xx or yy seen during model training. Zero-resource NMT (Currey and Heafield, 2019; Firat et al., 2016b; Bapna et al., 2022), on the other hand, generates pseudo-parallel data between direct language pairs, utilizing these synthetic xx-yy datasets during training. Yet bilingual or pivot-based models for direct NMT, though showing impressive NMT quality, are not practical and cost-efficient, especially as the number of languages for translation increases. Additionally, relying on related languages when evaluating a multilingual model on an unseen pair does not guarantee good NMT performance. At the same time, scaling up to multilingual models of billions of parameters, trained on very large datasets of multiple language pairs, induces large training and inference costs.

Sparse models, and specifically, sparsely activated Mixture-of-Experts (MoE) models

Work completed while interning at Google.




Figure 1: Task-level MoE model, with LP - based routing; each Language Pair is routed through a top-2 router to an expert in the model experts' layer. From a pretrained Task-level MoE model we can extract a smaller dense network specializing in a certain task, e.g. ja-ko NMT.

(Shazeer et al., 2017; Lepikhin et al., 2020; Zoph et al., 2022; Ryabinin and Gusev, 2020), shape a promising direction in creating a universal translation model between all languages, while addressing efficiency concerns. We are the first to investigate performance of a specific type of MoE models, Task-level MoE models (Kudugunta et al., 2021) in the context of Direct NMT, focusing on the impact of different task-to-expert mapping methods on model performance and expert utilization. In our experiments, we train large multilingual Task-level MoE models across various configurations, including Zero-shot, Zero-resource, and scenarios where direct language pair data is available. We conduct a comparative analysis of BLEU scores across Task-level MoE models with different mapping strategies and expert counts, as well as bilingual NMT and pivot-based NMT models; our findings serve as a valuable resource for determining the optimal model configuration to achieve high-quality translation results in various direct language pairs. By training large sparse multilingual Task-level MoE models which perform well for direct pairs, and showcase training efficiency, we may be able to extract smaller expert-specific dense models, which specialize in translating between Language Pairs or translating into one particular language. These models are ready for deployment and can replace bilingual or pivot-based models (Figure 1). Prior work on Task-level MoEs (Kudugunta et al., 2021) demonstrates the inference efficiency of task-level routing, compared to other approaches (token- or sentence-level), for the task of Multilingual NMT between English-centric pairs (En-yy and xx-En). To the best of our knowledge, we're the

first ones to examine performance of Task-level MoE models in Direct NMT, offering a practical solution for expanding the coverage and quality of NMT across diverse language pairs.

2 Related Work

Direct Neural Machine Translation

Firat et al. (2016b) first proposed a multilingual multi-way zero-resource NMT model which outperformed previous pivot-based approaches, laying the foundation for numerous subsequent zero-shot and/or zero-resource NMT methods. Various researchers have since expanded upon this foundational work, exploring different techniques and strategies for improving direct NMT (Ji et al., 2020; Johnson et al., 2017; Zhang et al., 2020; Chen et al., 2017; Lu et al., 2018; Al-Shedivat and Parikh, 2019; Rios et al., 2020; Zheng et al., 2017; Gu et al., 2019; Arivazhagan et al., 2019; Kim et al., 2019; Cheng and Cheng, 2019; Chen et al., 2018; Currey and Heafield, 2019; Ballesteros and Sanderson, 2003; Lakew et al., 2021). More specifically, Freitag and Firat (2020) constructed a multi-way aligned training set from existing training corpora, aligning examples in different language pairs with identical source or target sentences. They introduced a hierarchical data sampling strategy to prevent the over-representation of English in the training set and employed a temperature-based strategy to select target languages, while choosing source languages uniformly. Their complete Multilingual NMT models demonstrated improved performance on direct pairs and comparable to baseline results on English-centric pairs. Liu et al. (2020) discretized the encoder output latent space into entries in a learned codebook, enabling the

representation of source sentences in a common language and increasing model robustness and performance in zero-shot setups. More recently, ? pre-trained and fine-tuned a many-to-many multilingual NMT model that surpassed several pivoting and bilingual baselines for Direct NMT, and ElNokrashy et al. (2022), presented a novel approach in which source and target language tokens are appended to the encoder input, while target language tokens are added to the decoder input. This differs from previous methods that used Source-only tokens and beginning of Sentence tokens for the encoder and decoder inputs, respectively, and aims to improve performance of Direct NMT models by providing more relevant context to both the encoder and decoder.

Latest works include the approach of Yang et al. (2022), who unified three distinct models to guide the student model during Direct NMT, aiming to improve its performance by harnessing the strengths of multiple models; and Xu et al. (2022), who proposed a Direct NMT method comprising two stages: extraction of highly similar examples and their translations across different languages to construct a multi-way aligned parallel dataset, which enables the model to learn from a richer set of translation examples, and generation of additional aligned examples using a well-trained generative model.

Conditional Computation, Sparsity and Mixture-of-Experts models

Fedus et al. (2022) extensively review sparse Mixture-of-Experts (MoE) models in Deep Learning, starting from the foundational work of Shazeer et al. (2017), Lepikhin et al. (2020) and Fedus et al. (2021). Their work delves into the scaling characteristics, routing algorithms, and training improvements proposed for MoE models, as well as the latest applications of Sparse MoE models across various domains, such as Natural Language Processing (NLP) (Zoph et al., 2022; Chowdhery et al., 2022; Lee-Thorp and Ainslie, 2022; Du et al., 2022), Computer Vision (Riquelme et al., 2021; Puigcerver et al., 2020; Wu et al., 2022; Hwang et al., 2022), Speech Recognition (You et al., 2021, 2022), and Multimodal Learning (Mustafa et al., 2022). In NLP, numerous studies have explored the potential of Sparse MoE models for diverse applications. Kudugunta et al. (2021) introduced the concept of task-level routing in

MoE models, achieving improved translation scores compared to large multilingual models with the same inference cost and dense student models distilled from token- or position-level MoE models. By using task-level MoE models in our experiments, we investigate and optimize their performance in direct pairs' NMT.

3 Experiments

Model

We train sparse encoder-decoder Task-level MoE models, with dim = 1024, hidden dim = 4096, 8 heads, 3 layers. Following ?, we similarly replace the Transformer Feed Forward Network (FFN) with a set of identical FFN experts. We use a 32,000 token SentencePiece vocabulary (Kudo and Richardson, 2018), shared in both the encoder and decoder of the model. We experiment with the number of experts, selecting either 16 or 64 experts, which correspond to models with approximately 1 billion and 3.5 billion parameters, respectively. We investigate the impact of different task id to expert mapping methods, setting the task id to be either the Target Language or the Language Pair for translation. Following Kudo and Richardson (2018); Kudugunta et al. (2021), data sampling temperature is T=5 during training, and we train 8 models for 2 million steps.

The training process has a batch size of 256, a maximum sentence length of 128. By examining performance of these models under different configurations, we want to better understand the relationship between the number of experts, task id to experts' mappings, and translation quality, ultimately informing the optimal design choices for Task-level MoE models in various scenarios.

Languages and Training Datasets

We train our models using in-house production-scale datasets, consisting of English-centric, and Direct pairs' parallel sentences, covering 108 languages, including English. Our training data includes a total of 107 English-centric pairs, and 53 direct pairs, with data sizes for each language pair ranging from a few thousand to a few billion sentences for each pair. The number of sentences in each parallel pair in the train set is provided in Table 4 in the Appendix.

	ja-th		bg-mk		ru-tr		fr-ar	
	←	→	←	→	←	→	←	→
Task-level MoE	21.51	25.36	26.44	23.29	10.24	8.72	19.37	15.06
Bilingual	6.24	7.26	24.47	21.28	9.53	7.58	19.6	9.71
Pivot-level	18.34	39.28	28.24	27.03	18.44	13.84	21.57	14.91
Number of Experts	16	64	16	16	16	16	16	16
Decoding mapping method	lp_b	tl_b	tl_a	tl_a	tl_a	tl_a	tl_a	tl_a

	zh-ko		ja-vi		ja-ko		ja-zh	
	←	→	←	→	←	→	←	→
Task-level MoE	34.35	28.70	25.17	19.18	42.64	41.38	34.94	22.52
Bilingual	5.71	3.16	8.20	19.6	6.12	4.47	3.16	7.26
Pivot-level	28.13	35.24	18.97	19.61	21.49	36.69	19.98	27.46
Number of Experts	16	16	16	64	16	16	16	64
Decoding mapping method	lp_b	lp_b	lp_b	lp_a	lp_b	lp_b	lp_b	lp_a

Table 1: Best BLEU scores from our Task-level MoE models trained with 16 or 64 experts, for direct pairs. We compare to Bilingual and Pivot-level NMT models’ scores and state the Task-level MoE model/setup under which best scores were obtained. overall best scores per pair are highlighted. lp_x decoding mapping method implies a model trained with Language Pair (LP) - based task to expert mapping, and tl_x implies target language (TL) - based task to expert mapping.

	ja-th		bg-mk		ru-tr		fr-ar	
	←	→	←	→	←	→	←	→
Task-level MoE (LP)	lp_a	0.47 2.26	3.21 22.75	1.55 0.74	-	-	-	-
	lp_b	21.51 18.92	3.07 20.43	4.86 0.83	19.37	15.06		
	lp_c	0.42 1.50	5.97 18.32	5.57 0.80	2.40	0.86		
Task-level MoE (TL)	tl_a	1.15 0.78	26.44 23.29	10.24 8.72	5.36	1.18		
	tl_b	1.39 6.09	5.19 10.15	4.45 0.86	5.62	0.81		
Bilingual		6.24 7.26	24.47 21.28	9.53 7.58	19.60	9.71		
Pivot-level		18.34 39.28	28.24 27.03	18.44 13.84	21.57	14.91		
Number of train sentences		6,200	205,651	1,467,160		17,163,359		

	zh-ko		ja-vi		ja-ko		ja-zh	
	←	→	←	→	←	→	←	→
Task-level MoE (LP)	lp_a	2.81 4.09	0.77 5.32	3.33 11.78	1.67	9.06		
	lp_b	34.35 28.70	25.17 1.78	42.64 41.38	34.93	2.50		
	lp_c	2.35 3.04	0.57 1.31	2.71 10.14	1.30	2.26		
Task-level MoE (TL)	tl_a	4.51 0.61	0.87 14.31	3.87 1.30	1.86	18.77		
	tl_b	4.44 10.7	1.34 4.80	3.91 19.91	2.79	4.06		
Bilingual		5.71 3.16	8.20 19.6	6.12 4.47	3.16	7.26		
Pivot-level		28.13 35.24	18.97 19.61	21.49 36.69	19.98	27.46		
Number of train sentences		498,968	604,940	974,896		1,339,622		

Table 2: BLEU scores of Task-level MoE models trained with 16 experts, for direct pairs, in both directions, for models with Language Pair (LP) - or target language (TL) - based routing during training, lp_a, lp_b, lp_c and tl_a, tl_b mapping during inference, respectively. We compare to Bilingual and Pivot-level NMT models’ scores and highlight the model/setup under which best scores were obtained.

	ja-th		bg-mk	fr-ar		zh-ko		ja-vi	ja-ko		ja-zh
	←	→	→	←	→	←	→	←	←	→	←
Task-level MoE (LP)	lp_a	16.33 17.50	19.18	-	-	21.34 19.08	19.18	25.3	25.25	22.52	
	lp_b	0.44 2.79	19.07	4.47	2.13	2.05 2.51	0.47	0.96	4.46	0.71	
	lp_c	0.36 0.50	19.04	2.69	1.18	0.59 0.47	0.45	0.80	2.05	0.56	
Task-level MoE (TL)	tl_a	8.30 25.33	20.29	7.37	5.38	12.61 20.27	5.95	14.25	26.39	0.01	
	tl_b	8.33 25.36	20.35	7.32	5.50	12.43 20.34	5.96	14.26	26.40	10.49	
Bilingual		6.24 7.26	21.28	9.71	19.60	5.71 3.16	8.20	6.12	4.47	3.16	
Pivot-level		18.34 39.28	27.03	14.91	21.57	28.13 35.24	18.97	21.49	36.69	19.98	
Number of train sentences		6,200 205,651			17,163,359	498,968		604,940		974,896	1,339,622

Table 3: BLEU scores of Task-level MoE models trained with 64 experts, for direct pairs, in both directions, for models with Language Pair (LP) - or target language (TL) - based routing during training, lp_a, lp_b, lp_c and tl_a, tl_b mapping during inference, respectively. We compare to Bilingual and Pivot-level NMT models’ scores and highlight the model/setup under which best scores were obtained.

Evaluation Sets and Task Mapping

For evaluation, we use internal test sets from two distinct sources: the Web Domain and Wikipedia. These test sets vary in size, containing anywhere from 500 to 5,000 sentences for each language pair. In both training and evaluation datasets, we follow a specific preprocessing convention, prepending the source sentence with `<4xx>` `<2yy>` tokens. These tokens serve as a cue for the model, helping it identify the source and target languages in context, and the model to better learn and adapt to various language pair combinations.

Our models use either the exact train set Language Pair (LP) or the target language (TL) of the pair as the task id for mapping translation pairs to different experts. In the former case, we directly map the specific language pair (En-yy, xx-En or xx-yy) to a task id, and have a total of 214 tasks (107 x 2 English-centric pairs) when training with English-only datasets, or 267 tasks (+ 53 direct Language Pairs) when incorporating Direct LPs. On the other hand, when using the target language as the task id, we have 108 tasks (107 languages and English) in all scenarios. There we leverage and try to benefit from the shared characteristics of translations into the same target language. Since we evaluate our models on direct (non-English) data Language Pairs, it is often the case that the exact mapping of the pair to an expert was not defined during training. To address this issue, we experiment with several approaches for mapping language pair xx-yy to a task id at inference time - in each method mentioned below we name the component used as task id:

- Language Pair (LP) as task id.
 - **lp_mapping_a (lp_a)**: exact LP
 - **lp_mapping_b (lp_b)**: En-yy LP
 - **lp_mapping_c (lp_c)**: xx-En LP
- Target language as task id
 - **tl_mapping_a (tl_a)**: yy
 - **tl_mapping_b (tl_b)**: xx

4 Results & Discussion

In this study, we have two primary baselines for comparison: bilingual NMT models and pivot-based NMT models, which are trained using a bridging strategy through English. We use BLEU (Papineni et al., 2002) for evaluating

the translation quality of all direct language pairs. We are aware that learned metrics like COMET (Rei et al., 2022) and BLEURT (Sellam et al., 2020) show higher correlation with human judgements for high resource, English centric language pairs (Freitag et al., 2022). However, it is unknown how well learned metrics perform on low resource languages and thus we decide to stick with BLEU for the main part of the paper. Tables 2 and 3 summarize results for Task-level MoE models with either 16 or 64 experts.

In Table 1, we show the specific configurations that led to the highest BLEU scores for Task-level MoE models - comparing those to the performance of bilingual and pivot-based models, we aim to emphasize the best BLEU scores for each language pair. This analysis allows us to identify the most effective models for each pair and gain insights into how different model configurations and expert counts contribute to translation performance, so we can better understand the advantages and limitations of Task-level MoE models with varying numbers of experts, as well as bilingual and pivot-based NMT models. Our results ultimately help inform the choice of translation models and strategies to optimize performance for different language pairs

Performance of NMT models varies significantly depending on the pair, and on the direction of translation - certain pairs show large discrepancies between translation directions. We see that pivot-based models perform better than bilingual ones for the majority of pairs. Bilingual NMT and pivot-based NMT models demonstrate strong results in several instances, while Task-level MoE models show mixed performance across different Language Pairs.

Which model performs best for each pair?

Table 1 results reveal that for the majority of direct language pairs, the highest Task-level MoE BLEU scores are achieved using a model with 16 experts, with the `tl_a` and `lp_b` mapping-to-experts methods surpassing other approaches. It is particularly noteworthy that for pairs of languages belonging to closely related language families (Japanese, Chinese, Korean, Vietnamese, and Thai), the LP-mapping-based models seem to perform best. In contrast, TL-based models tend to outperform the rest of the other language pairs. In this case, Task-level MoE models show superior performance in




Figure 2: Difference between BLEU scores of Task-level MoE models trained with 16 and 64 experts, for direct pairs, for models with Language Pair (LP) - based routing during training, lp_a, lp_b, lp_c mapping during inference, respectively, and models with Target Language (TL) - based routing during training, tl_a, tl_b mapping during inference. For each pair, we mark the BLEU score value of the model which shows the largest difference among the 16 and 64 experts Task-level MoE variation.

the forward translation direction across all pairs. Meanwhile, pivot-based models yield the best results in the backward translation direction for four out of five pairs in those related languages. This observation suggests that different mapping-to-experts strategies may be more effective for certain language pairs or translation directions, depending on the linguistic relationships and unique characteristics of the languages involved; a tailored approach in employing Task-level MoE models and mapping-to-experts methods is the key to optimizing translation performance for different language pairs and directions.

Comparison of different task to expert mapping methods

From Table 2, pivot-based models, Task-level MoE (LP) models with lp_b task id mapping and Task-level MoE (TL) models with tl_a task id mapping (both trained with 16 experts) are primarily the best performing models among all pairs.

Task-level MoE models with LP task id to expert routing and lp_b mapping at inference time excel in translations to Japanese, from Thai, Vietnamese, Korean, and Chinese (+3.17, +6.2, +21.15, +14.95 BLEU points over the pivot-level NMT model, respectively, and +15.27, +16.97, +36.52, +31.77 BLEU gain against bilingual models) as well as Korean-Chinese (ko-zh) (+6.22 BLEU compared to pivot-based model, +28.64 BLEU over the bilingual baseline) and French-

Arabic (fr-ar) (+0.15 and +5.35 BLEU over the pivot and bilingual NMT models, respectively). This indicates the lp_b mapping strategy is particularly effective for these language pairs; routing to the expert best at En-yy translation, where yy is each pair’s target language, instead of routing to the expert explicitly trained on the exact direct pair, xx-yy, likely captures better the linguistic nuances of the pair and leads to improved translation scores.

Task-level MoE models trained with the TL task-to-expert mapping approach and the tl_a task id mapping, which leverages target language pairs for routing during inference, achieve the highest scores for translations involving Kazakh (kk), such as Kazakh-Russian (kk-ru) in both directions (+13.62 and +13.68 BLEU gain against pivot-based model, +0.34 and +12.34 BLEU points over bilingual models). Success of tl_a mapping strategy suggests the model shows superior performance in those cases and when the pair is properly routed to the expert specialised in the pair’s target language.

Table 3 shows results from our evaluation of 64 experts’ Task-level MoE models. Interestingly, pivot-based methods stand out as top performing in the majority of pairs and directions. These pairs include Bulgarian-Macedonian (bg-mk), French-Arabic (fr-ar), and Japanese-Thai (ja-th), as well as Chinese-Korean (zh-ko) in both directions, and Japanese-Korean (ja-ko). In all other pairs and directions, specifically when translating to

Japanese from Vietnamese, Korean, and Chinese, Task-level MoE (LP) models with the `lp_a` mapping approach emerge as the best-performing models (+0.21, +3.81, +2.54 BLEU points against the pivot-based model, respectively, and +10, +19.18, +19.36 BLEU gain over the bilingual baseline).

For certain Language Pairs with smaller amounts of training data bilingual NMT models achieve reasonable performance even with limited data and often outperform others. On the other hand, we see that pivot-based NMT models perform well when there is a substantial amount of training data (e.g., Belarusian–Macedonian, and Russian–Turkish in both directions, Arabic–French, Chinese–Korean, Japanese–{Thai, Vietnamese, Chinese}). This might be due to the fact that these models leverage English as an intermediate bridging language between source and target languages, which is helpful when there’s enough data to help in learning meaningful representations.

Comparison of models with different number of experts

In Figure 2, we present the difference in BLEU scores on select direct language pairs for Task-level MoE models trained using Language Pair (LP) or Target Language (TL) based mapping, with 16 and 64 experts. We seek to visually emphasize differences in performance across models that utilize varying numbers of experts, in order to get insights into their relative effectiveness and understand how the choice of mapping strategy (LP or TL) and the number of experts (16 or 64) impacts translation performance. For the majority of pairs (Vietnamese–Japanese, and both directions of French–Arabic, Japanese–Thai, Chinese–Korean, Japanese–Korean, Japanese–Chinese), Task-level MoE models with 16 experts outperform those with 64. For those language pairs, with the exception of Japanese–Chinese, for which TL-based Task-level MoE with `tl_a` shows best performance, LP-based Task-level MoE with `lp_b` mapping show the largest gain when trained with fewer experts, corroborating our previous findings. For other pairs, such as Japanese–Thai, Japanese–Chinese, and Bengali–Macedonian, different model variations show a largest improvement when trained with either 16 or 64 experts. We also notice a significant, yet not maximum, BLEU score gain for a large number of pairs with the TL-based Task-level MoE model and the LP-based Task-level MoE model, when using `tl_a` and `lp_a` on the task-

to-expert mapping during inference, respectively, between 16 and 64 experts’ models.

Routing decisions analysis

In Figures 3, 4, 5, 6 we show the experts’ utilization for the last model’s encoder and decoder layers for different pairs, for our TL-based Task-level MoE model, with `tl_a` used as a task to expert mapping approach during training; The darker a cell is in each of the Figures, the more that expert is used by the specific language pair. We can observe that language pairs with the same target language get properly routed to the same expert, as expected. It’s interesting to see that there are no major differences between the encoder’s routing decisions across checkpoints, yet we notice changes in experts’ assignment in the decoder during and at the end of training; at first, pairs with the same target language are mapped to the same different experts, yet as the model converges, we see a significant overlap in the selected experts, for pairs with different target languages. We also see a great overlap in experts chosen in the encoder during and at the end of training - the majority of experts are the same throughout model training and certain experts are preferred over others. This preference differs from the experts chosen in the decoder, where we see almost no overlap. It is also worth noting that the number of common experts between the encoder and the decoder is minimal, both during training and upon model convergence.

In the Appendix, we additionally show experts’ utilization in the encoder and decoder of our 64 experts’ TL-based Task-level MoE model.

5 Conclusions

We conducted a thorough analysis of the usage of Task-level Mixture of Expert models for Direct NMT. Our experiments reveal the strengths and weaknesses of different approaches and shed light on which configurations work best for specific direct language pairs. Our comparisons help in identifying best-performing models and offer valuable insights into how varying numbers of experts, and different task-to-expert mapping methods, during training and inference, can influence direct pairs’ translation quality in Task-level MoE models. Specifically, we noticed that Task-level MoE NMT models, along with pivot-based approaches, are frequently top performers for numerous direct language pairs. However, their




Figure 3: Routing decisions of the last encoder layer of our Task-level MoE model with 16 experts, trained with pair target language to task id mapping, with tl_a used during inference, for 1M steps.




Figure 4: Routing decisions of the last decoder layer of our Task-level MoE model with 16 experts, trained with pair target language to task id mapping, with tl_a used during inference, for 1M steps.

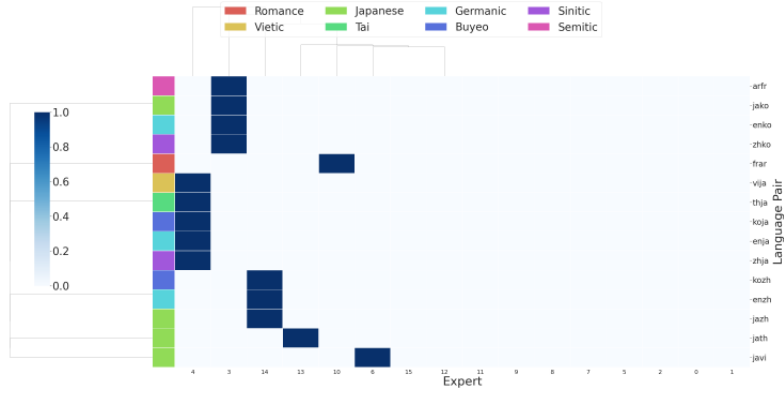


Figure 5: Routing decisions of the last encoder layer of our Task-level MoE model with 16 experts, trained with pair target language to task id mapping, with tl_a used during inference, for 2M steps.




Figure 6: Routing decisions of the last decoder layer of our Task-level MoE model with 16 experts, trained with pair target language to task id mapping, with tl_a used during inference, for 2M steps.

performance varies between different pairs and translation directions. Besides observing NMT quality results, our analysis of expert’s usage throughout training also serves as an informative way to visualize and understand mapping of language pairs to model experts in the encoder and decoder throughout training. Future work can focus on enhancing this model for broader language coverage, including other low-resource languages, to further improve translation quality and efficiency.

Limitations

Training and evaluation of Task-level MoE models can be very challenging due to the models’ size and complexity. There is a large number of parameters that require tuning, such as the number and size of experts, the number and diversity of languages, and of English-centric and direct Language Pairs in training, the batch size, the vocabulary size and the maximum sentence length. At the same time, it is time and computationally expensive to evaluate on a large number of direct pairs, train separate Bilingual and Pivot-level models for all to use as baselines, and also perform the experts’ analysis and visualizations for those pairs. This automatically restricts the amount and breadth of results, and calls for further exploration of the model capabilities in the future.

Ethics Statement

Working with large language models raises several ethical concerns in NLP research, particularly related to quality, bias and toxicity of the output result ([Bender et al., 2021](#); [Chowdhery et al., 2022](#); [Blodgett et al., 2020](#); [Brown et al., 2020](#)). In the context of NMT, involvement of many stakeholders in the task, as well as the dangers arising from mistranslation of the original text need to be taken into consideration, since they can potentially affect the perception of the work and harm the interests of the different parts ([Taivalkoski-Shilov, 2019](#); [Gambier and Van Doorslaer, 2016](#)). Undoubtedly, the benefits of responsibly developed and deployed NMT systems lie in making an author’s work more accessible, enabling the transfer of ideas to other audiences, and enhancing both the reader’s capabilities and the translator’s role and sense of contribution ([Besacier, 2014](#)), which are directions any work in the area should aim to focus on.

References

- Roei Aharoni, Melvin Johnson, and Orhan Firat. 2019. [Massively multilingual neural machine translation](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 3874–3884, Minneapolis, Minnesota. Association for Computational Linguistics.
- Maruan Al-Shedivat and Ankur P Parikh. 2019. [Consistency by agreement in zero-shot neural machine translation](#). *arXiv preprint arXiv:1904.02338*.
- Naveen Arivazhagan, Ankur Bapna, Orhan Firat, Roei Aharoni, Melvin Johnson, and Wolfgang Macherey. 2019. [The missing ingredient in zero-shot neural machine translation](#). *arXiv preprint arXiv:1903.07091*.
- Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. 2014. [Neural machine translation by jointly learning to align and translate](#). *arXiv preprint arXiv:1409.0473*.
- Lisa Ballesteros and Mark Sanderson. 2003. [Addressing the lack of direct translation resources for cross-language retrieval](#). In *Proceedings of the twelfth international conference on Information and knowledge management*, pages 147–152.
- Ankur Bapna, Isaac Caswell, Julia Kreutzer, Orhan Firat, Daan van Esch, Aditya Siddhant, Mengmeng Niu, Pallavi Baljekar, Xavier Garcia, Wolfgang Macherey, et al. 2022. [Building machine translation systems for the next thousand languages](#). *arXiv preprint arXiv:2205.03983*.
- Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. [On the dangers of stochastic parrots: Can language models be too big?](#) In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency, FAccT ’21*, page 610–623, New York, NY, USA. Association for Computing Machinery.
- Laurent Besacier. 2014. [Machine translation for literature: a pilot study \(traduction automatisée d’une oeuvre littéraire: une étude pilote\) \[in French\]](#). In *Proceedings of TALN 2014 (Volume 2: Short Papers)*, pages 389–394, Marseille, France. Association pour le Traitement Automatique des Langues.
- Su Lin Blodgett, Solon Barocas, Hal Daumé III, and Hanna Wallach. 2020. [Language \(technology\) is power: A critical survey of “bias” in NLP](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5454–5476, Online. Association for Computational Linguistics.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. [Language models are few-shot learners](#). *Advances in neural information processing systems*, 33:1877–1901.
- Yun Chen, Yang Liu, Yong Cheng, and Victor OK Li. 2017. [A teacher-student framework for zero-resource neural machine translation](#). *arXiv preprint arXiv:1705.00753*.
- Yun Chen, Yang Liu, and Victor Li. 2018. [Zero-resource neural machine translation with multi-agent communication game](#). In *Proceedings of the aaai conference on artificial intelligence*, volume 32.
- Yong Cheng and Yong Cheng. 2019. [Joint training for pivot-based neural machine translation](#). *Joint training for neural machine translation*, pages 41–54.
- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. 2022. [Palm: Scaling language modeling with pathways](#). *arXiv preprint arXiv:2204.02311*.
- Anna Currey and Kenneth Heafield. 2019. [Zero-resource neural machine translation with monolingual pivot data](#). In *Proceedings of the 3rd Workshop on Neural Generation and Translation*, pages 99–107.
- Nan Du, Yanping Huang, Andrew M Dai, Simon Tong, Dmitry Lepikhin, Yuanzhong Xu, Maxim Krikun, Yanqi Zhou, Adams Wei Yu, Orhan Firat, et al. 2022. [Glam: Efficient scaling of language models with mixture-of-experts](#). In *International Conference on Machine Learning*, pages 5547–5569. PMLR.
- Muhammad ElNokrashy, Amr Hendy, Mohamed Maher, Mohamed Afify, and Hany Hassan Awadalla. 2022. [Language tokens: A frustratingly simple approach improves zero-shot performance of multilingual translation](#). *arXiv preprint arXiv:2208.05852*.
- William Fedus, Jeff Dean, and Barret Zoph. 2022. [A review of sparse expert models in deep learning](#). *arXiv preprint arXiv:2209.01667*.
- William Fedus, Barret Zoph, and Noam Shazeer. 2021. [Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity](#). *J. Mach. Learn. Res.*, 23:1–40.
- Orhan Firat, Kyunghyun Cho, and Yoshua Bengio. 2016a. [Multi-way, multilingual neural machine translation with a shared attention mechanism](#). *arXiv preprint arXiv:1601.01073*.
- Orhan Firat, Baskaran Sankaran, Yaser Al-Onaizan, Fatos T Yarman Vural, and Kyunghyun Cho. 2016b. [Zero-resource translation with multilingual neural machine translation](#). *arXiv preprint arXiv:1606.04164*.

- Markus Freitag and Orhan Firat. 2020. [Complete multilingual neural machine translation](#). *arXiv preprint arXiv:2010.10239*.
- Markus Freitag, Ricardo Rei, Nitika Mathur, Chi-kiu Lo, Craig Stewart, Eleftherios Avramidis, Tom Kocmi, George Foster, Alon Lavie, and André F. T. Martins. 2022. [Results of WMT22 metrics shared task: Stop using BLEU – neural metrics are better and more robust](#). In *Proceedings of the Seventh Conference on Machine Translation (WMT)*, pages 46–68, Abu Dhabi, United Arab Emirates (Hybrid). Association for Computational Linguistics.
- Yves Gambier and Luc Van Doorslaer. 2016. *Border crossings: Translation studies and other disciplines*, volume 126. John Benjamins Publishing Company.
- Jiatao Gu, Yong Wang, Kyunghyun Cho, and Victor OK Li. 2019. [Improved zero-shot neural machine translation via ignoring spurious correlations](#). *arXiv preprint arXiv:1906.01181*.
- Hany Hassan, Anthony Aue, Chang Chen, Vishal Chowdhary, Jonathan Clark, Christian Federmann, Xuedong Huang, Marcin Junczys-Dowmunt, William Lewis, Mu Li, et al. 2018. [Achieving human parity on automatic chinese to english news translation](#). *arXiv preprint arXiv:1803.05567*.
- Changho Hwang, Wei Cui, Yifan Xiong, Ziyue Yang, Ze Liu, Han Hu, Zilong Wang, Rafael Salas, Jithin Jose, Prabhat Ram, et al. 2022. [TutEL: Adaptive mixture-of-experts at scale](#). *arXiv preprint arXiv:2206.03382*.
- Baijun Ji, Zhirui Zhang, Xiangyu Duan, Min Zhang, Boxing Chen, and Weihua Luo. 2020. [Cross-lingual pre-training based transfer for zero-shot neural machine translation](#). In *Proceedings of the AAAI conference on artificial intelligence*, volume 34, pages 115–122.
- Melvin Johnson, Mike Schuster, Quoc V Le, Maxim Krikun, Yonghui Wu, Zhifeng Chen, Nikhil Thorat, Fernanda Viégas, Martin Wattenberg, Greg Corrado, et al. 2017. [Google’s multilingual neural machine translation system: Enabling zero-shot translation](#). *Transactions of the Association for Computational Linguistics*, 5:339–351.
- Yunsu Kim, Petre Petrov, Pavel Petrushkov, Shahram Khadivi, and Hermann Ney. 2019. [Pivot-based transfer learning for neural machine translation between non-english languages](#). *arXiv preprint arXiv:1909.09524*.
- Taku Kudo and John Richardson. 2018. [Sentencepiece: A simple and language independent subword tokenizer and detokenizer for neural text processing](#). *arXiv preprint arXiv:1808.06226*.
- Sneha Kudugunta, Yanping Huang, Ankur Bapna, Maxim Krikun, Dmitry Lepikhin, Minh-Thang Luong, and Orhan Firat. 2021. [Beyond distillation: Task-level mixture-of-experts for efficient inference](#). *arXiv preprint arXiv:2110.03742*.
- Surafel M Lakew, Matteo Negri, and Marco Turchi. 2021. [Zero-shot neural machine translation with self-learning cycle](#). In *Proceedings of the 4th Workshop on Technologies for MT of Low Resource Languages (LoResMT2021)*, pages 96–113.
- James Lee-Thorp and Joshua Ainslie. 2022. [Sparse mixers: Combining moe and mixing to build a more efficient bert](#). *arXiv preprint arXiv:2205.12399*.
- Dmitry Lepikhin, Hyoungho Lee, Yuanzhong Xu, Dehao Chen, Orhan Firat, Yanping Huang, Maxim Krikun, Noam Shazeer, and Zhifeng Chen. 2020. [Gshard: Scaling giant models with conditional computation and automatic sharding](#). *arXiv preprint arXiv:2006.16668*.
- Danni Liu, Jan Niehues, James Cross, Francisco Guzmán, and Xian Li. 2020. [Improving zero-shot translation by disentangling positional information](#). *arXiv preprint arXiv:2012.15127*.
- Yichao Lu, Phillip Keung, Faisal Ladhak, Vikas Bhardwaj, Shaonan Zhang, and Jason Sun. 2018. [A neural interlingua for multilingual machine translation](#). *arXiv preprint arXiv:1804.08198*.
- Basil Mustafa, Carlos Riquelme, Joan Puigcerver, Rodolphe Jenatton, and Neil Houlsby. 2022. [Multimodal contrastive learning with limoe: the language-image mixture of experts](#).
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. [Bleu: a method for automatic evaluation of machine translation](#). In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 311–318, Philadelphia, Pennsylvania, USA. Association for Computational Linguistics.
- Jerin Philip, Alexandre Berard, Matthias Gallé, and Laurent Besacier. 2020. [Monolingual adapters for zero-shot neural machine translation](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4465–4470, Online. Association for Computational Linguistics.
- Joan Puigcerver, Carlos Riquelme, Basil Mustafa, Cédric Renggli, André Susano Pinto, Sylvain Gelly, Daniel Keysers, and Neil Houlsby. 2020. [Scalable transfer learning with expert models](#). *CoRR*, abs/2009.13239.
- Ricardo Rei, José G. C. de Souza, Duarte Alves, Chrysoula Zerva, Ana C Farinha, Taisiya Glushkova, Alon Lavie, Luisa Coheur, and André F. T. Martins. 2022. [COMET-22: Unbabel-IST 2022 submission for the metrics shared task](#). In *Proceedings of the Seventh Conference on Machine Translation (WMT)*, pages 578–585, Abu Dhabi, United Arab Emirates (Hybrid). Association for Computational Linguistics.
- Annette Rios, Mathias Müller, and Rico Sennrich. 2020. [Subword segmentation and a single bridge language affect zero-shot neural machine translation](#). *arXiv preprint arXiv:2011.01703*.

- Carlos Riquelme, Joan Puigcerver, Basil Mustafa, Maxim Neumann, Rodolphe Jenatton, André Susano Pinto, Daniel Keysers, and Neil Houlsby. 2021. [Scaling vision with sparse mixture of experts](#). *Advances in Neural Information Processing Systems*, 34:8583–8595.
- Max Ryabinin and Anton Gusev. 2020. [Towards crowdsourced training of large neural networks using decentralized mixture-of-experts](#).
- Thibault Sellam, Dipanjan Das, and Ankur Parikh. 2020. [BLEURT: Learning robust metrics for text generation](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 7881–7892, Online. Association for Computational Linguistics.
- Noam Shazeer, Azalia Mirhoseini, Krzysztof Maziarczyk, Andy Davis, Quoc Le, Geoffrey Hinton, and Jeff Dean. 2017. [Outrageously large neural networks: The sparsely-gated mixture-of-experts layer](#). *arXiv preprint arXiv:1701.06538*.
- Ilya Sutskever, Oriol Vinyals, and Quoc V Le. 2014. [Sequence to sequence learning with neural networks](#). *Advances in neural information processing systems*, 27.
- Kristiina Taivalkoski-Shilov. 2019. [Ethical issues regarding machine \(-assisted\) translation of literary texts](#). *Perspectives*, 27(5):689–703.
- Yuqing Tang, Chau Tran, Xian Li, Peng-Jen Chen, Naman Goyal, Vishrav Chaudhary, Jiatao Gu, and Angela Fan. 2021. [Multilingual translation from denoising pre-training](#). In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, pages 3450–3466, Online. Association for Computational Linguistics.
- Jörg Tiedemann. 2018. [Emerging language spaces learned from massively multilingual corpora](#). *arXiv preprint arXiv:1802.00273*.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. [Attention is all you need](#). *Advances in neural information processing systems*, 30.
- Weizhi Wang, Zhirui Zhang, Yichao Du, Boxing Chen, Jun Xie, and Weihua Luo. 2021. [Rethinking zero-shot neural machine translation: From a perspective of latent variables](#). In *Findings of the Association for Computational Linguistics: EMNLP 2021*, pages 4321–4327, Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Lemeng Wu, Mengchen Liu, Yinpeng Chen, Dongdong Chen, Xiyang Dai, and Lu Yuan. 2022. [Residual mixture of experts](#). *arXiv preprint arXiv:2204.09636*.
- Yulin Xu, Zhen Yang, Fandong Meng, et al. 2022. [Eag: Extract and generate multi-way aligned corpus for complete multi-lingual neural machine translation](#). *arXiv preprint arXiv:2203.02180*.
- Jian Yang, Yuwei Yin, Shuming Ma, Dongdong Zhang, Shuangzhi Wu, Hongcheng Guo, Zhoujun Li, and Furu Wei. 2022. [UM4: Unified multilingual multiple teacher-student model for zero-resource neural machine translation](#). In *Proceedings of the Thirty-First International Joint Conference on Artificial Intelligence*. International Joint Conferences on Artificial Intelligence Organization.
- Yilin Yang, Akiko Eriguchi, Alexandre Muzio, Prasad Tadepalli, Stefan Lee, and Hany Hassan. 2021. [Improving multilingual translation by representation and gradient regularization](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7266–7279, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Zhao You, Shulin Feng, Dan Su, and Dong Yu. 2021. [Speechmoe: Scaling to large acoustic models with dynamic routing mixture of experts](#). *CoRR*, abs/2105.03036.
- Zhao You, Shulin Feng, Dan Su, and Dong Yu. 2022. [Speechmoe2: Mixture-of-experts model with improved routing](#). In *ICASSP 2022-2022 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 7217–7221. IEEE.
- Biao Zhang, Philip Williams, Ivan Titov, and Rico Sennrich. 2020. [Improving massively multilingual neural machine translation and zero-shot translation](#). *arXiv preprint arXiv:2004.11867*.
- Hao Zheng, Yong Cheng, and Yang Liu. 2017. [Maximum expected likelihood estimation for zero-resource neural machine translation](#). In *IJCAI*, pages 4251–4257.
- Barret Zoph, Irwan Bello, Sameer Kumar, Nan Du, Yanping Huang, Jeff Dean, Noam Shazeer, and William Fedus. 2022. [St-moe: Designing stable and transferable sparse expert models](#). *arXiv preprint arXiv:2202.08906*.
- Barret Zoph and Kevin Knight. 2016. [Multi-source neural translation](#). In *Proceedings of the 2016 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, pages 30–34, San Diego, California. Association for Computational Linguistics.

Appendix

In Figures 7, 8 we see the routing decisions maps for the last encoder and decoder layer of TL- based task-level MoE models trained with 64 experts and `tl_a` mapping during inference. The distribution to the experts is similarly expected here, following the training and inference pairs-to-experts' mapping, as same target language pairs get routed to the same experts. The overlap in the encoder and decoder experts is minimal. All but one selected experts are different between the encoder and decoder.




Figure 7: Routing decisions of the last encoder layer of our Task-level MoE model with 64 experts, trained with pair target language to task id mapping, with tl_a used during inference, for 2M steps.




Figure 8: Routing decisions of the last decoder layer of our Task-level MoE model with 64 experts, trained with pair target language to task id mapping, with tl_a used during inference, for 2M steps.

LP	Num. of sentences	LP	Num. of sentences	LP	Num. of sentences
ar-fr	0	fr-ar	0	be-ru	3,641
ru-be	3,641	ja-th	6,200	th-ja	6,200
ko-th	6,208	th-zh	12,305	zh-th	12,305
id-ms	24,535	kk-tr	69,476	tr-kk	69,476
ug-en	129,688	tk-en	140,247	en-tk	144,094
en-ug	167,741	kk-ru	177,931	ru-kk	177,931
bg-mk	205,651	mk-bg	205,651	or-en	262,099
cs-es	285,546	es-cs	285,546	ko-vi	369,594
vi-ko	369,594	et-fi	466,254	fi-et	466,254
ko-zh	498,968	zh-ko	498,968	vi-zh	504,037
zh-vi	504,037	en-tt	556,751	de-ru	585,161
ru-de	585,161	ja-vi	604,940	vi-ja	604,940
zh-tr	609,066	en-or	658,015	en-rw	792,559
tt-en	819,579	cs-fr	895,912	fr-cs	895,912
ja-ko	974,896	ko-ja	974,896	cs-de	1,007,448
de-cs	1,007,448	rw-en	1,264,813	ja-zh	1,339,622
zh-ja	1,339,622	ru-tr	1,467,160	tr-ru	1,467,160
cs-ru	1,644,795	ru-cs	1,644,795	en-la	1,809,003
en-xh	1,895,974	la-en	2,167,039	en-sm	2,353,608
de-es	2,886,854	de-fr	2,886,854	es-de	2,886,854
en-st	2,931,225	en-sn	3,022,722	en-ig	3,058,540
en-haw	3,393,044	en-yo	3,404,254	en-lo	3,573,485
fr-de	3,681,694	ru-uk	3,792,978	uk-ru	3,792,978
en-yi	3,904,022	en-ku	5,254,260	en-ny	5,408,239
en-so	5,472,253	xh-en	5,728,449	sm-en	5,730,016
lo-en	5,766,691	en-mi	5,793,624	en-co	6,847,178
ig-en	7,224,372	en-tg	8,136,182	sn-en	8,319,539
yi-en	8,347,482	en-zu	8,376,755	haw-en	8,646,782
en-ha	8,995,196	en-sd	9,095,582	en-ps	9,446,484
yo-en	9,565,981	st-en	10,052,460	en-ky	10,213,875
mi-en	10,383,802	en-ht	11,758,532	ku-en	11,984,479
en-ga	12,494,184	ny-en	12,616,421	so-en	12,745,144
en-su	14,088,043	en-am	14,583,643	en-mt	15,693,771
en-pa	15,784,955	en-ceb	15,965,215	ga-en	16,145,438
en-mg	17,474,049	co-en	17,490,449	en-be	17,645,900
en-eo	18,104,746	es-ru	18,363,087	ru-es	18,363,087
tg-en	18,509,169	en-fy	19,495,193	ha-en	19,687,379
en-mn	20,012,644	sd-en	20,337,227	es-fr	20,792,781
fr-es	20,792,781	zu-en	21,010,498	ky-en	21,357,639
ps-en	21,613,753	ht-en	22,455,117	en-gd	22,495,036
fr-ru	22,890,960	ru-fr	22,890,960	en-eu	23,279,562
en-my	23,848,386	en-hy	25,779,512	en-lb	25,943,545
be-en	28,585,822	en-hmn	29,258,771	pa-en	29,592,573
su-en	30,234,526	en-jv	30,315,732	en-uz	30,943,271
en-si	31,847,809	en-kk	32,075,645	en-cy	32,089,017
am-en	32,110,524	en-ml	32,335,094	mg-en	32,356,401
en-gu	32,742,711	fy-en	33,389,347	en-mk	33,546,951
en-mr	33,629,504	hy-en	34,146,553	eu-en	35,158,779
hmn-en	35,795,085	en-bs	35,801,135	gd-en	36,082,889
en-kn	36,472,587	en-km	37,229,149	en-hr	37,248,537
en-ne	38,606,074	en-sw	41,198,888	mt-en	41,723,852

LP	Num. of sentences	LP	Num. of sentences	LP	Num. of sentences
en-ka	42,992,730	en-te	43,112,561	en-gl	45,053,738
lb-en	46,042,055	kn-en	46,847,179	en-sr	48,786,193
mn-en	49,510,283	cy-en	50,252,795	en-af	50,525,196
gu-en	53,217,408	en-sq	55,374,368	my-en	55,886,059
lv-en	57,400,218	ceb-en	59,055,133	mk-en	61,867,654
en-ta	62,352,671	eo-en	62,552,455	en-az	64,274,455
kk-en	64,724,987	ne-en	69,288,148	si-en	69,667,030
en-lv	69,998,174	ml-en	70,059,552	en-is	71,837,680
hr-en	72,301,373	en-ur	74,947,168	en-et	75,581,372
te-en	75,967,139	en-bn	80,362,347	km-en	80,514,558
en-uk	86,997,037	en-ca	87,267,893	en-sk	89,304,972
gl-en	89,695,691	bs-en	90,263,002	sw-en	90,493,834
af-en	94,492,215	en-fil	97,193,343	en-sl	97,924,827
ka-en	99,344,373	ta-en	100,659,244	en-lt	102,081,855
is-en	103,425,627	sq-en	104,573,595	mr-en	105,247,834
uz-en	105,701,543	az-en	115,304,502	en-iw	117,411,687
en-ms	117,562,320	en-fa	117,720,531	en-bg	128,893,255
en-fi	133,307,056	en-el	135,547,629	en-ro	136,723,562
ca-en	137,014,332	sr-en	137,913,360	en-hu	146,210,128
ur-en	149,852,215	fil-en	156,986,036	lv-en	163,533,602
en-no	165,068,512	en-cs	167,802,683	en-da	176,676,171
sl-en	189,776,212	bn-en	190,650,787	et-en	197,134,423
uk-en	209,888,151	en-sv	217,128,777	lt-en	245,092,644
bg-en	289,282,624	sk-en	295,623,418	fi-en	300,151,301
el-en	322,429,692	th-en	340,378,951	ro-en	352,986,741
en-th	353,001,311	fa-en	402,836,484	ms-en	414,693,796
iw-en	427,057,571	hu-en	436,082,952	hi-en	490,464,288
en-hi	494,638,634	ar-en	525,411,304	id-en	527,345,119
en-ar	532,778,336	ko-en	535,863,357	en-id	538,581,029
da-en	539,573,001	en-ko	546,191,335	no-en	613,221,604
vi-en	632,740,475	en-vi	656,115,977	cs-en	661,833,636
sv-en	800,882,060	ja-en	846,991,020	tr-en	850,154,583
en-tr	869,295,018	en-ja	876,842,917	it-en	998,195,505
en-it	100,981,5294	pl-en	111,935,3071	en-pl	1,141,598,628
pt-en	1,155,038,272	en-pt	1,184,401,180	en-zh	1,228,817,744
zh-en	1,238,691,743	ru-en	1,425,268,039	en-ru	1,455,103,126
nl-en	1,580,819,532	en-nl	1,587,530,791	de-en	1,680,270,443
en-de	1,695,095,726	en-fr	1,887,609,530	fr-en	1,922,463,803
en-es	2,419,825,975	es-en	2,435,228,645		

Table 4: Number of Sentences of each Language Pair (LP) in our train set.